

# Course 02224

## Scheduling

### Basic Notions

Hans Henrik Løvengreen

DTU Compute, Technical University of Denmark

## General Scheduling Notions

- Given
  - ▶ A number of *activities* (tasks, jobs) to be done  
Each may be *one-instance* or *recurrent*
  - ▶ A number of *resources* to be utilized
  - ▶ A specification of the resources needed for each task
- A *schedule* is a mapping over time between tasks and resources consistent with the specification.
- *Scheduling* is the process of determining a schedule
- Scheduling is done by a *scheduler* according to a *policy*
- Scheduling may be *static* (planning) or *dynamic*

## Scheduling Objectives

### Activity Types

- *Batch* activities — High CPU/IO ratio
- *Interactive* activities — Low CPU/IO ratio
- *Real-time* activities — deadlines, timeliness

### Goals

- *Fairness* — every process/user/client get a fair share
- Good *utilization* — no resource is unnecessarily idle
- High *throughput*
- Acceptable *response times*
- *Deadlines* must be met

## Real-Time Scheduling Problem

### Real-time Programs

- A *real-time program* is a reactive program with real-time requirements.
- Real-time requirements may be *hard*, *firm*, or *soft*.

### Real-Time Scheduling

- Given a real-time program, the *scheduling problem* is to ensure that the (hard) real-time requirements are met given a limited number of HW resources
- Traditionally, mostly CPU-resources have been considered.
- Other resources may also be considered:  
Communication bandwidth, power consumption ...

## Real-Time Scheduling Notions

### Schedulability

- Real-time requirements can usually be stated as *deadlines*
- A *feasible schedule* is a schedule which meets the deadlines

### Scheme

- A *scheduling scheme* consists of:
  - ▶ A *scheduling policy* determining the possible schedules
  - ▶ A *schedulability test* ensuring that all schedules are feasible
- A test may be *sufficient* and/or *necessary*
- A test is *sustainable* if improving conditions preserves feasibility

### Optimality

- A scheduling policy is *optimal* (for a given class of scheduling problems) iff it can find a feasible schedule, whenever one exists

## The Task Execution Model

An arbitrary program is too complex to analyse.

To facilitate analysis, an *abstract execution model* is used:

- The program consists of a fixed number  $N$  of computational *tasks*
- Tasks are recurrent: **loop**
  - await* next release
  - compute task
  - respond
- Tasks are *released* periodically with fixed *periods*  $T_i$
- Tasks have fixed *worst-case execution times*  $C_i$
- Tasks have fixed response *deadlines*  $D_i$  ( $C_i \leq D_i \leq T_i$ )
- Tasks may have fixed *start times* (offsets)  $S_i$  ( $S_i \geq 0$ )

## The Simple Scheduling Problem

- Tasks have no suspension points
- *Overheads* (context switch, scheduling, ... ) are ignored
- Only a execution on a *single CPU* resource is considered
- Tasks are assumed to be *independent*:
  - ▶ No sharing of resources other resources than CPU(s)
  - ▶ No precedence relation
  - ▶ No assumption on start times
- No aperiodic tasks
- Tasks must run within their period:  $D_i = T_i$

## Static Scheduling

### Idea


- A feasible, repeatable schedule is determined statically.
- The schedule is executed by a single *cyclic executive* driven by a regular hardware tick.

### Properties

- + Simple implementation not requiring an operating system
- Very difficult to adapt to task modifications
- Cannot utilize spare time for other activities
- Not amendable for multi-processing

## Cyclic Executive

Task	$T$	$C$
a	25	10
b	25	8
c	50	5
d	50	4
e	100	2




## Thread Based Task Scheduling

- Tasks are performed repeatedly by dedicated threads scheduled by an OS.
- A task may be *ready* or *waiting*.
- A *scheduler* dynamically chooses a ready task for execution
- General policies:
  - ▶ Fair scheduling, e.g. *round robin* time-sharing
  - ▶ Priority based scheduling
- Fair scheduling not suited for hard real-time
- Priorities may be *fixed* or *dynamic*.

## Round-Robin Scheduling

- Ready tasks are executed in *time slots* in a cyclic way


**Example** time slot = 2.0



## Round-Robin Scheduling

- Ready threads are executed in *time slots* in a cyclic way

**Example** time slot = 0.5




## Fixed Priority Scheduling (FPS)

- **Idea:** Assign fixed priorities to tasks
- Scheduler needs to know about priorities (only)
- Always choose highest priority task when rescheduling
- Scheduling points:
  - ▶ *Non-preemptive scheduling:* Let current task run till completion
  - ▶ *Periodic scheduling:* Schedule at periodic intervals
  - ▶ *Deferred scheduling:* Schedule after a given time
  - ▶ *Preemptive scheduling:* Preempt current task if higher prioritized task becomes ready
- How should equal priority tasks be treated?
- What is the best way of assigning priorities?


## Fixed Priority Scheduling

### Example



## Fixed Priority Scheduling

### Example



## Rate-Monotonic Priority Assignment

- A rate-monotonic priority assignment (RMA) satisfies:

$$T_i < T_j \Rightarrow P_i > P_j$$

- Optimal for the simple scheduling problem:


*If a set of tasks is schedulable for some fixed priority assignment, then:*

*The set of tasks is schedulable using rate-monotonic priority assignment.*



## Fixed Priority Scheduling

### RMA Example



- Feasible if all deadlines met from critical instant till longest period

## Utilization

- The processor *utilization* (load) contributed by task  $i$ :

$$U_i = \frac{C_i}{T_i}$$

- For a set of tasks running on a single CPU:

$$\sum_{i=1}^N \frac{C_i}{T_i} \leq 1$$

- How high can we go while ensuring schedulability?

## Utilization based check

In 1973, Liu and Layland showed a sufficient condition for schedulability (using rate-monotonic priority assignment):

$$\sum_{i=1}^N \frac{C_i}{T_i} < N(2^{1/N} - 1)$$

- For various  $N$  this gives:

$N$	Utilization (%)
1	100.0
2	82.8
3	78.0
4	75.7
10	71.8
$\infty$	69.3 ( $\ln 2$ )

- Bound sometimes too pessimistic
- Refinement:  $N$  = number of *task families* (having harmonic periods)

## Utilization based check

In 2006, Bini et.al. showed an alternative sufficient condition:

$$\prod_{i=1}^N \left( \frac{C_i}{T_i} + 1 \right) \leq 2$$

### Example

- Consider again:
- | Task | $T$ | $C$ |
|------|-----|-----|
| a    | 3   | 1   |
| b    | 4   | 1   |
| c    | 5   | 1   |

- $(1/3 + 1) \cdot (1/4 + 1) \cdot (1/5 + 1) = (4 \cdot 5 \cdot 6) / (3 \cdot 4 \cdot 5) = 2$  ✓

## Response Time Analysis (RTA)

### Idea

- Assume some priority assignment given.
- Calculate worst-case *response time*  $R_i$  for each task.
- If  $R_i \leq T_i$  for all  $i$ , the tasks are schedulable.

The response time satisfies:

$$R_i^{k+1} = C_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i^k}{T_j} \right\rceil C_j$$

- $R_i$  can be found by an iterative process.

## Response Time Analysis

### Example

	$T$	$C$	$R^0$	$R^1$	$R^2$	$R^3$	...
a	50	10	10	<u>10</u>			
b	70	15	15	25	<u>25</u>		
c	110	40	40	65	75	90	<u>90</u>
d	200	30	30	95	120	170	195 <u>195</u>

## Meeting Deadlines

- Some tasks may require  $D_i < T_i$
- If  $R_i \leq D_i$  for all processes, all deadlines are met
- A *deadline monotonic* priority assignment satisfies:

$$D_i < D_j \Rightarrow P_i > P_j$$

- Deadline monotonic priority assignment is optimal for independent processes

### Proof (sketch)

- Given a feasible priority assignment  $W$
- Swap any two adjacent tasks with  $P_i > P_j$  but  $D_i > D_j$
- When none left, assignment is feasible and deadline monotonic
- Corollary: RMA is optimal for  $D_i = T_i$

## Sporadic Tasks

- *Aperiodic tasks* are released at arbitrary times by external or internal events
- *Sporadic tasks* have a *minimum arrival time*  $T_i$
- Sporadic tasks can be treated as periodic tasks with period  $T_i$
- Alternatively, aperiodic and sporadic tasks may be handled by a *server task* given a certain fraction of the processing time
- Sporadic tasks often have  $D_i < T_i$  (eg. *alarms*)

## Dynamic Priority Scheduling

- **Idea:** Base scheduling choice on current situation
- Scheduler needs to be aware of more task parameters
- Better exploitation of processor time

### Common Schemes

- Least Completion Time (LCT)
- Least Slack time (LST)
- Earliest Deadline First (EDF)

## Earliest Deadline First (EDF)

### Principle

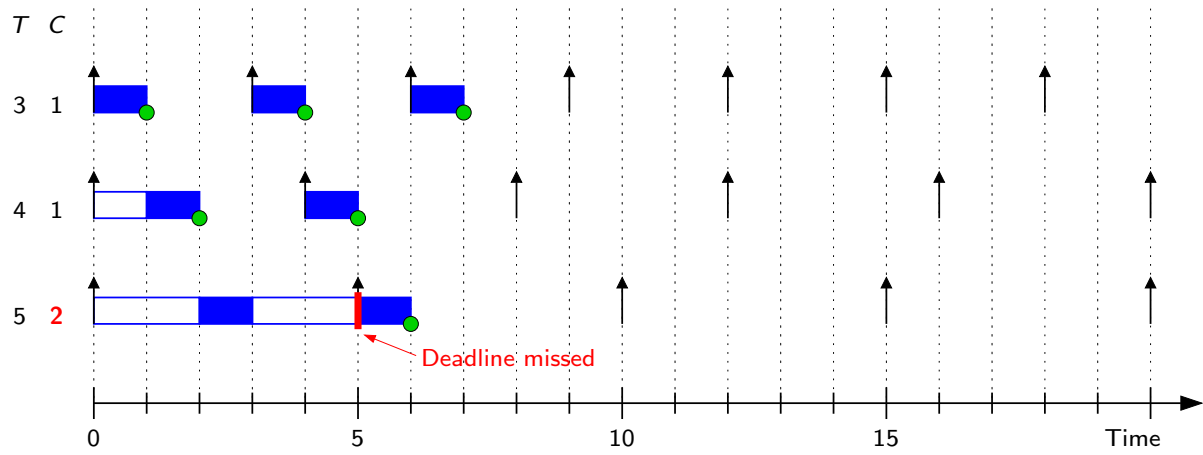
- Always run tasks with first-coming deadline

### Properties

- + Optimal algorithm: Feasible iff  $\sum U_i \leq 1$  (for  $D_i = T_i$ )
- Scheduler must know about and maintain deadlines
- Performs poorly on overload
- Difficult to analyse response times
- Difficult to analyse for non-preemptive processes

## EDF Scheduling

### Example FPS



## EDF Scheduling

### Example EDF

